Jane Doe

jane.doe@email.com | (555) 123-4567 | linkedin.com/in/janedoe | github.com/janedoe | San Francisco, CA

SKILLS

Languages: Python, SQL, R, Scala, Java

ML/AI: PyTorch, TensorFlow, Scikit-learn, XGBoost, LightGBM, Keras, Statistical Learning, Feature Engineering

MLOps & Cloud: AWS (SageMaker, EC2, S3, Lambda), Docker, Kubernetes, MLflow, Weights & Biases, Apache Airflow

Data & Analytics: Spark, Pandas, NumPy, PostgreSQL, MongoDB, Redis, Jupyter, Model Optimization, A/B Testing

EXPERIENCE

Senior Machine Learning Engineer | TechCorp Inc | Jan 2021 - Present

• Deployed 12+ production ML models using PyTorch and AWS SageMaker, serving 5M+ daily predictions with 99.9% uptime

• Optimized deep learning model inference latency by 60% through feature engineering and model compression techniques

• Built end-to-end MLOps pipeline using Docker and Kubernetes, reducing model deployment time from 2 weeks to 2 hours

• Led statistical learning initiatives that improved recommendation system CTR by 35%, generating $2M additional revenue

• Designed automated feature engineering framework processing 50TB+ daily data, increasing model accuracy by 18%

Machine Learning Engineer | DataTech Solutions | Mar 2019 - Dec 2020

• Developed production ML pipelines using TensorFlow and Scikit-learn, processing 10M+ transactions daily

• Implemented model optimization strategies reducing training time by 45% while maintaining 92% prediction accuracy

• Created real-time feature store using Redis and Apache Spark, supporting 20+ ML models across 3 product teams

• Collaborated with data science team to deploy ensemble models achieving 25% improvement in customer churn prediction

• Established ML monitoring system using MLflow, tracking model performance across 15+ production environments

Junior ML Engineer | StartupAI | Jun 2018 - Feb 2019

• Built recommendation engine using collaborative filtering and PyTorch, increasing user engagement by 40%

• Performed statistical analysis and A/B testing on ML models, validating performance improvements with 95% confidence

• Engineered features from unstructured text data using NLP techniques, improving sentiment classification F1-score to 0.89

• Optimized hyperparameters using Bayesian optimization, reducing model training cost by 30% on AWS infrastructure

PROJECTS

Real-time Fraud Detection System | Led development of ensemble ML system using XGBoost and TensorFlow, achieving 94% precision and processing 100K+ transactions per minute

Large Language Model Fine-tuning | Fine-tuned transformer models for domain-specific tasks, achieving 15% performance improvement over baseline using parameter-efficient methods

EDUCATION

Master of Science, Computer Science | Stanford University | 2018

Bachelor of Science, Mathematics | UC Berkeley | 2016

CERTIFICATIONS

AWS Certified Machine Learning - Specialty | 2022
Google Professional Machine Learning Engineer | 2021